

18.3. `select` — Waiting for I/O completion

This module provides access to the `select()` and `poll()` functions available in most operating systems, `devpoll()` available on Solaris and derivatives, `epoll()` available on Linux 2.5+ and `kqueue()` available on most BSD. Note that on Windows, it only works for sockets; on other operating systems, it also works for other file types (in particular, on Unix, it works on pipes). It cannot be used on regular files to determine whether a file has grown since it was last read.

Note: The `selectors` module allows high-level and efficient I/O multiplexing, built upon the `select` module primitives. Users are encouraged to use the `selectors` module instead, unless they want precise control over the OS-level primitives used.

The module defines the following:

exception `select.error`

A deprecated alias of `OSError`.

Changed in version 3.3: Following [PEP 3151](#), this class was made an alias of `OSError`.

`select.devpoll()`

(Only supported on Solaris and derivatives.) Returns a `/dev/poll` polling object; see section [/dev/poll Polling Objects](#) below for the methods supported by `devpoll` objects.

`devpoll()` objects are linked to the number of file descriptors allowed at the time of instantiation. If your program reduces this value, `devpoll()` will fail. If your program increases this value, `devpoll()` may return an incomplete list of active file descriptors.

The new file descriptor is [non-inheritable](#).

New in version 3.3.

Changed in version 3.4: The new file descriptor is now non-inheritable.

`select.epoll(sizehint=-1, flags=0)`

(Only supported on Linux 2.5.44 and newer.) Return an edge polling object, which can be used as Edge or Level Triggered interface for I/O events. *sizehint* and *flags* are deprecated and completely ignored.

See the [Edge and Level Trigger Polling \(epoll\) Objects](#) section below for the methods supported by epolling objects.

epoll objects support the context management protocol: when used in a [with](#) statement, the new file descriptor is automatically closed at the end of the block.

The new file descriptor is [non-inheritable](#).

Changed in version 3.3: Added the *flags* parameter.

Changed in version 3.4: Support for the [with](#) statement was added. The new file descriptor is now non-inheritable.

Deprecated since version 3.4: The *flags* parameter. `select.EPOLL_CLOEXEC` is used by default now. Use [os.set_inheritable\(\)](#) to make the file descriptor inheritable.

`select.poll()`

(Not supported by all operating systems.) Returns a polling object, which supports registering and unregistering file descriptors, and then polling them for I/O events; see section [Polling Objects](#) below for the methods supported by polling objects.

`select.kqueue()`

(Only supported on BSD.) Returns a kernel queue object; see section [Kqueue Objects](#) below for the methods supported by kqueue objects.

The new file descriptor is [non-inheritable](#).

Changed in version 3.4: The new file descriptor is now non-inheritable.

`select.kevent(ident, filter=KQ_FILTER_READ, flags=KQ_EV_ADD, fflags=0, data=0, udata=0)`

(Only supported on BSD.) Returns a kernel event object; see section [Kevent Objects](#) below for the methods supported by kevent objects.

`select.select(rlist, wlist, xlist[, timeout])`

This is a straightforward interface to the Unix `select()` system call. The first three arguments are sequences of ‘waitable objects’: either integers representing file descriptors or objects with a parameterless method named [fileno\(\)](#) returning such an integer:

- *rlist*: wait until ready for reading
- *wlist*: wait until ready for writing
- *xlist*: wait for an “exceptional condition” (see the manual page for what your system considers such a condition)

Empty sequences are allowed, but acceptance of three empty sequences is platform-dependent. (It is known to work on Unix but not on Windows.) The optional *timeout* argument specifies a time-out as a floating point number in seconds. When the *timeout* argument is omitted the function blocks until at least one file descriptor is ready. A time-out value of zero specifies a poll and never blocks.

The return value is a triple of lists of objects that are ready: subsets of the first three arguments. When the time-out is reached without a file descriptor becoming ready, three empty lists are returned.

Among the acceptable object types in the sequences are Python [file objects](#) (e.g. `sys.stdin`, or objects returned by `open()` or `os.popen()`), socket objects returned by `socket.socket()`. You may also define a *wrapper* class yourself, as long as it has an appropriate `fileno()` method (that really returns a file descriptor, not just a random integer).

Note: File objects on Windows are not acceptable, but sockets are. On Windows, the underlying `select()` function is provided by the WinSock library, and does not handle file descriptors that don’t originate from WinSock.

Changed in version 3.5: The function is now retried with a recomputed timeout when interrupted by a signal, except if the signal handler raises an exception (see [PEP 475](#) for the rationale), instead of raising `InterruptedError`.

`select.PIPE_BUF`

The minimum number of bytes which can be written without blocking to a pipe when the pipe has been reported as ready for writing by `select()`, `poll()` or another interface in this module. This doesn’t apply to other kind of file-like objects such as sockets.

This value is guaranteed by POSIX to be at least 512. Availability: Unix.

New in version 3.2.

18.3.1. /dev/poll Polling Objects

Solaris and derivatives have `/dev/poll`. While `select()` is $O(\text{highest file descriptor})$ and `poll()` is $O(\text{number of file descriptors})$, `/dev/poll` is $O(\text{active file descriptors})$.

`/dev/poll` behaviour is very close to the standard `poll()` object.

`devpoll.close()`

Close the file descriptor of the polling object.

New in version 3.4.

`devpoll.closed`

True if the polling object is closed.

New in version 3.4.

`devpoll.fileno()`

Return the file descriptor number of the polling object.

New in version 3.4.

`devpoll.register(fd[, eventmask])`

Register a file descriptor with the polling object. Future calls to the `poll()` method will then check whether the file descriptor has any pending I/O events. `fd` can be either an integer, or an object with a `fileno()` method that returns an integer. File objects implement `fileno()`, so they can also be used as the argument.

`eventmask` is an optional bitmask describing the type of events you want to check for. The constants are the same that with `poll()` object. The default value is a combination of the constants `POLLIN`, `POLLPRI`, and `POLLOUT`.

Warning: Registering a file descriptor that's already registered is not an error, but the result is undefined. The appropriate action is to unregister or modify it first. This is an important difference compared with `poll()`.

`devpoll.modify(fd[, eventmask])`

This method does an `unregister()` followed by a `register()`. It is (a bit) more efficient than doing the same explicitly.

`devpoll.unregister(fd)`

Remove a file descriptor being tracked by a polling object. Just like the `register()` method, `fd` can be an integer or an object with a `fileno()` method that returns an integer.

Attempting to remove a file descriptor that was never registered is safely ignored.

`devpoll.poll([timeout])`

Polls the set of registered file descriptors, and returns a possibly-empty list containing (`fd`, `event`) 2-tuples for the descriptors that have events or errors to report. `fd` is the file descriptor, and `event` is a bitmask with bits set for the reported events for that descriptor — `POLLIN` for waiting input, `POLLOUT` to indicate that the descriptor can be written to, and so forth. An empty list indicates that the call timed out and no file descriptors had any events to report. If `timeout` is given, it specifies the length of time in milliseconds which the system will wait for events before returning. If `timeout` is omitted, `-1`, or `None`, the call will block until there is an event for this poll object.

Changed in version 3.5: The function is now retried with a recomputed timeout when interrupted by a signal, except if the signal handler raises an exception (see [PEP 475](#) for the rationale), instead of raising `InterruptedError`.

18.3.2. Edge and Level Trigger Polling (epoll) Objects

<http://linux.die.net/man/4/epoll>

eventmask

| Constant | Meaning |
|-----------------------------|---|
| <code>EPOLLIN</code> | Available for read |
| <code>EPOLLOUT</code> | Available for write |
| <code>EPOLLPRI</code> | Urgent data for read |
| <code>EPOLLERR</code> | Error condition happened on the assoc. fd |
| <code>EPOLLHUP</code> | Hang up happened on the assoc. fd |
| <code>EPOLLET</code> | Set Edge Trigger behavior, the default is Level Trigger behavior |
| <code>EPOLLONESHOT</code> | Set one-shot behavior. After one event is pulled out, the fd is internally disabled |
| <code>EPOLLEXCLUSIVE</code> | |

| Constant | Meaning |
|-------------|---|
| | Wake only one epoll object when the associated fd has an event. The default (if this flag is not set) is to wake all epoll objects polling on a fd. |
| EPOLLRDHUP | Stream socket peer closed connection or shut down writing half of connection. |
| EPOLLRDNORM | Equivalent to EPOLLIN |
| EPOLLRDBAND | Priority data band can be read. |
| EPOLLWRNORM | Equivalent to EPOLLOUT |
| EPOLLWRBAND | Priority data may be written. |
| EPOLLMMSG | Ignored. |

epoll.close()

Close the control file descriptor of the epoll object.

epoll.closed

True if the epoll object is closed.

epoll.fileno()

Return the file descriptor number of the control fd.

epoll.fromfd(fd)

Create an epoll object from a given file descriptor.

epoll.register(fd[, eventmask])

Register a fd descriptor with the epoll object.

epoll.modify(fd, eventmask)

Modify a registered file descriptor.

epoll.unregister(fd)

Remove a registered file descriptor from the epoll object.

epoll.poll(timeout=-1, maxevents=-1)

Wait for events. timeout in seconds (float)

Changed in version 3.5: The function is now retried with a recomputed timeout when interrupted by a signal, except if the signal handler raises an exception (see [PEP 475](#) for the rationale), instead of raising [InterruptedError](#).

18.3.3. Polling Objects

The `poll()` system call, supported on most Unix systems, provides better scalability for network servers that service many, many clients at the same time. `poll()` scales better because the system call only requires listing the file descriptors of interest, while `select()` builds a bitmap, turns on bits for the fds of interest, and then afterward the whole bitmap has to be linearly scanned again. `select()` is $O(\text{highest file descriptor})$, while `poll()` is $O(\text{number of file descriptors})$.

`poll.register(fd[, eventmask])`

Register a file descriptor with the polling object. Future calls to the `poll()` method will then check whether the file descriptor has any pending I/O events. `fd` can be either an integer, or an object with a `fileno()` method that returns an integer. File objects implement `fileno()`, so they can also be used as the argument.

`eventmask` is an optional bitmask describing the type of events you want to check for, and can be a combination of the constants `POLLIN`, `POLLPRI`, and `POLLOUT`, described in the table below. If not specified, the default value used will check for all 3 types of events.

| Constant | Meaning |
|------------------------|---|
| <code>POLLIN</code> | There is data to read |
| <code>POLLPRI</code> | There is urgent data to read |
| <code>POLLOUT</code> | Ready for output: writing will not block |
| <code>POLLERR</code> | Error condition of some sort |
| <code>POLLHUP</code> | Hung up |
| <code>POLLRDHUP</code> | Stream socket peer closed connection, or shut down writing half of connection |
| <code>POLLNVAL</code> | Invalid request: descriptor not open |

Registering a file descriptor that's already registered is not an error, and has the same effect as registering the descriptor exactly once.

`poll.modify(fd, eventmask)`

Modifies an already registered `fd`. This has the same effect as `register(fd, eventmask)`. Attempting to modify a file descriptor that was never registered causes an `OSError` exception with `errno ENOENT` to be raised.

`poll.unregister(fd)`

Remove a file descriptor being tracked by a polling object. Just like the `register()` method, `fd` can be an integer or an object with a `fileno()` method that returns an integer.

Attempting to remove a file descriptor that was never registered causes a `KeyError` exception to be raised.

`poll.poll([timeout])`

Polls the set of registered file descriptors, and returns a possibly-empty list containing `(fd, event)` 2-tuples for the descriptors that have events or errors to report. `fd` is the file descriptor, and `event` is a bitmask with bits set for the reported events for that descriptor — `POLLIN` for waiting input, `POLLOUT` to indicate that the descriptor can be written to, and so forth. An empty list indicates that the call timed out and no file descriptors had any events to report. If `timeout` is given, it specifies the length of time in milliseconds which the system will wait for events before returning. If `timeout` is omitted, negative, or `None`, the call will block until there is an event for this poll object.

Changed in version 3.5: The function is now retried with a recomputed timeout when interrupted by a signal, except if the signal handler raises an exception (see [PEP 475](#) for the rationale), instead of raising `InterruptedError`.

18.3.4. Kqueue Objects

`kqueue.close()`

Close the control file descriptor of the kqueue object.

`kqueue.closed`

True if the kqueue object is closed.

`kqueue.fileno()`

Return the file descriptor number of the control fd.

`kqueue.fromfd(fd)`

Create a kqueue object from a given file descriptor.

`kqueue.control(changelist, max_events[, timeout=None])` → eventlist

Low level interface to kevent

- `changelist` must be an iterable of kevent object or `None`
- `max_events` must be 0 or a positive integer
- `timeout` in seconds (floats possible)

Changed in version 3.5: The function is now retried with a recomputed timeout when interrupted by a signal, except if the signal handler raises an exception (see [PEP 475](#) for the rationale), instead of raising `InterruptedError`.

18.3.5. Kevent Objects

<https://www.freebsd.org/cgi/man.cgi?query=kqueue&sektion=2>

`kevent.ident`

Value used to identify the event. The interpretation depends on the filter but it's usually the file descriptor. In the constructor `ident` can either be an int or an object with a `fileno()` method. `kevent` stores the integer internally.

`kevent.filter`

Name of the kernel filter.

| Constant | Meaning |
|-------------------------------|---|
| <code>KQ_FILTER_READ</code> | Takes a descriptor and returns whenever there is data available to read |
| <code>KQ_FILTER_WRITE</code> | Takes a descriptor and returns whenever there is data available to write |
| <code>KQ_FILTER_AIO</code> | AIO requests |
| <code>KQ_FILTER_VNODE</code> | Returns when one or more of the requested events watched in <i>fflag</i> occurs |
| <code>KQ_FILTER_PROC</code> | Watch for events on a process id |
| <code>KQ_FILTER_NETDEV</code> | Watch for events on a network device [not available on Mac OS X] |
| <code>KQ_FILTER_SIGNAL</code> | Returns whenever the watched signal is delivered to the process |
| <code>KQ_FILTER_TIMER</code> | Establishes an arbitrary timer |

`kevent.flags`

Filter action.

| Constant | Meaning |
|---------------------------|--|
| <code>KQ_EV_ADD</code> | Adds or modifies an event |
| <code>KQ_EV_DELETE</code> | Removes an event from the queue |
| <code>KQ_EV_ENABLE</code> | Permits <code>control()</code> to return the event |
| | |

| Constant | Meaning |
|----------------|---|
| KQ_EV_DISABLE | Disablesevent |
| KQ_EV_ONESHOT | Removes event after first occurrence |
| KQ_EV_CLEAR | Reset the state after an event is retrieved |
| KQ_EV_SYSFLAGS | internal event |
| KQ_EV_FLAG1 | internal event |
| KQ_EV_EOF | Filter specific EOF condition |
| KQ_EV_ERROR | See return values |

kevent.**fflags**

Filter specific flags.

KQ_FILTER_READ and KQ_FILTER_WRITE filter flags:

| Constant | Meaning |
|---------------|-----------------------------------|
| KQ_NOTE_LOWAT | low water mark of a socket buffer |

KQ_FILTER_VNODE filter flags:

| Constant | Meaning |
|----------------|--------------------------------|
| KQ_NOTE_DELETE | <i>unlink()</i> was called |
| KQ_NOTE_WRITE | a write occurred |
| KQ_NOTE_EXTEND | the file was extended |
| KQ_NOTE_ATTRIB | an attribute was changed |
| KQ_NOTE_LINK | the link count has changed |
| KQ_NOTE_RENAME | the file was renamed |
| KQ_NOTE_REVOKE | access to the file was revoked |

KQ_FILTER_PROC filter flags:

| Constant | Meaning |
|-------------------|--|
| KQ_NOTE_EXIT | the process has exited |
| KQ_NOTE_FORK | the process has called <i>fork()</i> |
| KQ_NOTE_EXEC | the process has executed a new process |
| KQ_NOTE_PCTRLMASK | internal filter flag |
| KQ_NOTE_PDATAMASK | internal filter flag |

| Constant | Meaning |
|------------------|---|
| KQ_NOTE_TRACK | follow a process across <i>fork()</i> |
| KQ_NOTE_CHILD | returned on the child process for <i>NOTE_TRACK</i> |
| KQ_NOTE_TRACKERR | unable to attach to a child |

KQ_FILTER_NETDEV filter flags (not available on Mac OS X):

| Constant | Meaning |
|------------------|-----------------------|
| KQ_NOTE_LINKUP | link is up |
| KQ_NOTE_LINKDOWN | link is down |
| KQ_NOTE_LINKINV | link state is invalid |

kevent.**data**

Filter specific data.

kevent.**udata**

User defined value.