

## 21.12. `http.client` — HTTP protocol client

**Source code:** <Lib/http/client.py>

This module defines classes which implement the client side of the HTTP and HTTPS protocols. It is normally not used directly — the module `urllib.request` uses it to handle URLs that use HTTP and HTTPS.

**See also:** The [Requests package](#) is recommended for a higher-level HTTP client interface.

**Note:** HTTPS support is only available if Python was compiled with SSL support (through the `ssl` module).

The module provides the following classes:

`class http.client.HTTPConnection(host, port=None, [timeout, ] source_address=None)`

An `HTTPConnection` instance represents one transaction with an HTTP server. It should be instantiated passing it a host and optional port number. If no port number is passed, the port is extracted from the host string if it has the form `host:port`, else the default HTTP port (80) is used. If the optional `timeout` parameter is given, blocking operations (like connection attempts) will timeout after that many seconds (if it is not given, the global default timeout setting is used). The optional `source_address` parameter may be a tuple of a (host, port) to use as the source address the HTTP connection is made from.

For example, the following calls all create instances that connect to the server at the same host and port:

```
>>> h1 = http.client.HTTPConnection('www.python.org')
>>> h2 = http.client.HTTPConnection('www.python.org:80')
>>> h3 = http.client.HTTPConnection('www.python.org', 80)
>>> h4 = http.client.HTTPConnection('www.python.org', 80, timeout=
```

*Changed in version 3.2:* `source_address` was added.

*Changed in version 3.4:* The `strict` parameter was removed. HTTP 0.9-style “Simple Responses” are not longer supported.

```
class http.client.HTTPSConnection(host, port=None, key_file=None,
cert_file=None, [timeout, ]source_address=None, *, context=None,
check_hostname=None)
```

A subclass of [HTTPConnection](#) that uses SSL for communication with secure servers. Default port is 443. If *context* is specified, it must be a [ssl.SSLContext](#) instance describing the various SSL options.

Please read [Security considerations](#) for more information on best practices.

*Changed in version 3.2:* *source\_address*, *context* and *check\_hostname* were added.

*Changed in version 3.2:* This class now supports HTTPS virtual hosts if possible (that is, if [ssl.HAS\\_SNI](#) is true).

*Changed in version 3.4:* The *strict* parameter was removed. HTTP 0.9-style “Simple Responses” are no longer supported.

*Changed in version 3.4.3:* This class now performs all the necessary certificate and hostname checks by default. To revert to the previous, unverified, behavior [ssl.\\_create\\_unverified\\_context\(\)](#) can be passed to the *context* parameter.

*Deprecated since version 3.6:* *key\_file* and *cert\_file* are deprecated in favor of *context*. Please use [ssl.SSLContext.load\\_cert\\_chain\(\)](#) instead, or let [ssl.create\\_default\\_context\(\)](#) select the system’s trusted CA certificates for you.

The *check\_hostname* parameter is also deprecated; the [ssl.SSLContext.check\\_hostname](#) attribute of *context* should be used instead.

```
class http.client.HTTPResponse(sock, debuglevel=0, method=None,
url=None)
```

Class whose instances are returned upon successful connection. Not instantiated directly by user.

*Changed in version 3.4:* The *strict* parameter was removed. HTTP 0.9 style “Simple Responses” are no longer supported.

The following exceptions are raised as appropriate:

*exception* `http.client.HTTPException`

The base class of the other exceptions in this module. It is a subclass of [Exception](#).

*exception* `http.client.NotConnected`

A subclass of [HTTPException](#).

*exception* `http.client.InvalidURL`

A subclass of [HTTPException](#), raised if a port is given and is either non-numeric or empty.

*exception* `http.client.UnknownProtocol`

A subclass of [HTTPException](#).

*exception* `http.client.UnknownTransferEncoding`

A subclass of [HTTPException](#).

*exception* `http.client.UnimplementedFileMode`

A subclass of [HTTPException](#).

*exception* `http.client.IncompleteRead`

A subclass of [HTTPException](#).

*exception* `http.client.ImproperConnectionState`

A subclass of [HTTPException](#).

*exception* `http.client.CannotSendRequest`

A subclass of [ImproperConnectionState](#).

*exception* `http.client.CannotSendHeader`

A subclass of [ImproperConnectionState](#).

*exception* `http.client.ResponseNotReady`

A subclass of [ImproperConnectionState](#).

*exception* `http.client.BadStatusLine`

A subclass of [HTTPException](#). Raised if a server responds with a HTTP status code that we don't understand.

*exception* `http.client.LineTooLong`

A subclass of [HTTPException](#). Raised if an excessively long line is received in the HTTP protocol from the server.

*exception* `http.client.RemoteDisconnected`

A subclass of [ConnectionResetError](#) and [BadStatusLine](#). Raised by [HTTPConnection.getresponse\(\)](#) when the attempt to read the response results in no data read from the connection, indicating that the remote end has closed the connection.

*New in version 3.5:* Previously, [BadStatusLine\(''\)](#) was raised.

The constants defined in this module are:

`http.client.HTTP_PORT`

The default port for the HTTP protocol (always 80).

`http.client.HTTPS_PORT`

The default port for the HTTPS protocol (always 443).

`http.client.responses`

This dictionary maps the HTTP 1.1 status codes to the W3C names.

Example: `http.client.responses[http.client.NOT_FOUND]` is 'Not Found'.

See [HTTP status codes](#) for a list of HTTP status codes that are available in this module as constants.

## 21.12.1. HTTPConnection Objects

[HTTPConnection](#) instances have the following methods:

`HTTPConnection.request(method, url, body=None, headers={}, *, encode_chunked=False)`

This will send a request to the server using the HTTP request method *method* and the selector *url*.

If *body* is specified, the specified data is sent after the headers are finished. It may be a [str](#), a [bytes-like object](#), an open [file object](#), or an iterable of [bytes](#). If *body* is a string, it is encoded as ISO-8859-1, the default for HTTP. If it is a bytes-like object, the bytes are sent as is. If it is a [file object](#), the contents of the file is sent; this file object should support at least the `read()` method. If the file object is an instance of [io.TextIOBase](#), the data returned by the `read()` method will be encoded as ISO-8859-1, otherwise the data returned by `read()` is sent as is. If *body* is an iterable, the elements of the iterable are sent as is until the iterable is exhausted.

The *headers* argument should be a mapping of extra HTTP headers to send with the request.

If *headers* contains neither Content-Length nor Transfer-Encoding, but there is a request body, one of those header fields will be added automatically. If *body* is None, the Content-Length header is set to 0 for methods that expect a body (PUT, POST, and PATCH). If *body* is a string or a bytes-like object that is not also a [file](#), the Content-Length header is set to its length. Any other type of *body* (files and iterables in general) will be chunk-encoded, and the Transfer-Encoding header will automatically be set instead of Content-Length.

The `encode_chunked` argument is only relevant if Transfer-Encoding is specified in `headers`. If `encode_chunked` is `False`, the `HTTPConnection` object assumes that all encoding is handled by the calling code. If it is `True`, the body will be chunk-encoded.

**Note:** Chunked transfer encoding has been added to the HTTP protocol version 1.1. Unless the HTTP server is known to handle HTTP 1.1, the caller must either specify the Content-Length, or must pass a `str` or bytes-like object that is not also a file as the body representation.

*New in version 3.2:* `body` can now be an iterable.

*Changed in version 3.6:* If neither Content-Length nor Transfer-Encoding are set in `headers`, file and iterable `body` objects are now chunk-encoded. The `encode_chunked` argument was added. No attempt is made to determine the Content-Length for file objects.

### `HTTPConnection.getresponse()`

Should be called after a request is sent to get the response from the server. Returns an `HTTPResponse` instance.

**Note:** Note that you must have read the whole response before you can send a new request to the server.

*Changed in version 3.5:* If a `ConnectionError` or subclass is raised, the `HTTPConnection` object will be ready to reconnect when a new request is sent.

### `HTTPConnection.set_debuglevel(level)`

Set the debugging level. The default debug level is 0, meaning no debugging output is printed. Any value greater than 0 will cause all currently defined debug output to be printed to stdout. The `debuglevel` is passed to any new `HTTPResponse` objects that are created.

*New in version 3.1.*

### `HTTPConnection.set_tunnel(host, port=None, headers=None)`

Set the host and the port for HTTP Connect Tunnelling. This allows running the connection through a proxy server.

The host and port arguments specify the endpoint of the tunneled connection (i.e. the address included in the CONNECT request, *not* the address of the proxy server).

The `headers` argument should be a mapping of extra HTTP headers to send with the CONNECT request.

For example, to tunnel through a HTTPS proxy server running locally on port 8080, we would pass the address of the proxy to the `HTTPSConnection` constructor, and the address of the host that we eventually want to reach to the `set_tunnel()` method:

```
>>> import http.client
>>> conn = http.client.HTTPSConnection("localhost", 8080)
>>> conn.set_tunnel("www.python.org")
>>> conn.request("HEAD", "/index.html")
```

*New in version 3.2.*

#### `HTTPSConnection.connect()`

Connect to the server specified when the object was created. By default, this is called automatically when making a request if the client does not already have a connection.

#### `HTTPSConnection.close()`

Close the connection to the server.

As an alternative to using the `request()` method described above, you can also send your request step by step, by using the four functions below.

#### `HTTPSConnection.putrequest(method, url, skip_host=False, skip_accept_encoding=False)`

This should be the first call after the connection to the server has been made. It sends a line to the server consisting of the *method* string, the *url* string, and the HTTP version (HTTP/1.1). To disable automatic sending of Host: or Accept-Encoding: headers (for example to accept additional content encodings), specify *skip\_host* or *skip\_accept\_encoding* with non-False values.

#### `HTTPSConnection.putheader(header, argument[, ...])`

Send an [RFC 822](#)-style header to the server. It sends a line to the server consisting of the header, a colon and a space, and the first argument. If more arguments are given, continuation lines are sent, each consisting of a tab and an argument.

#### `HTTPSConnection.endheaders(message_body=None, *, encode_chunked=False)`

Send a blank line to the server, signalling the end of the headers. The optional *message\_body* argument can be used to pass a message body associated with the request.

If `encode_chunked` is `True`, the result of each iteration of `message_body` will be chunk-encoded as specified in [RFC 7230](#), Section 3.3.1. How the data is encoded is dependent on the type of `message_body`. If `message_body` implements the [buffer interface](#) the encoding will result in a single chunk. If `message_body` is a `collections.Iterable`, each iteration of `message_body` will result in a chunk. If `message_body` is a [file object](#), each call to `.read()` will result in a chunk. The method automatically signals the end of the chunk-encoded data immediately after `message_body`.

**Note:** Due to the chunked encoding specification, empty chunks yielded by an iterator body will be ignored by the chunk-encoder. This is to avoid premature termination of the read of the request by the target server due to malformed encoding.

*New in version 3.6:* Chunked encoding support. The `encode_chunked` parameter was added.

`HTTPConnection.send(data)`

Send data to the server. This should be used directly only after the [endheaders\(\)](#) method has been called and before [getresponse\(\)](#) is called.

## 21.12.2. HTTPResponse Objects

An [HTTPResponse](#) instance wraps the HTTP response from the server. It provides access to the request headers and the entity body. The response is an iterable object and can be used in a `with` statement.

*Changed in version 3.5:* The [io.BufferedIOBase](#) interface is now implemented and all of its reader operations are supported.

`HTTPResponse.read([amt])`

Reads and returns the response body, or up to the next *amt* bytes.

`HTTPResponse.readinto(b)`

Reads up to the next `len(b)` bytes of the response body into the buffer *b*. Returns the number of bytes read.

*New in version 3.3.*

`HTTPResponse.getheader(name, default=None)`

Return the value of the header *name*, or *default* if there is no header matching *name*. If there is more than one header with the name *name*, return all of the values joined by `' '`. If `'default'` is any iterable other than a single string, its elements are similarly returned joined by commas.

HTTPResponse.**getheaders()**

Return a list of (header, value) tuples.

HTTPResponse.**fileno()**

Return the fileno of the underlying socket.

HTTPResponse.**msg**

A `http.client.HTTPMessage` instance containing the response headers.  
`http.client.HTTPMessage` is a subclass of [email.message.Message](#).

HTTPResponse.**version**

HTTP protocol version used by server. 10 for HTTP/1.0, 11 for HTTP/1.1.

HTTPResponse.**status**

Status code returned by server.

HTTPResponse.**reason**

Reason phrase returned by server.

HTTPResponse.**debuglevel**

A debugging hook. If [debuglevel](#) is greater than zero, messages will be printed to stdout as the response is read and parsed.

HTTPResponse.**closed**

Is True if the stream is closed.

## 21.12.3. Examples

Here is an example session that uses the GET method:

```
>>> import http.client
>>> conn = http.client.HTTPSConnection("www.python.org")
>>> conn.request("GET", "/")
>>> r1 = conn.getresponse()
>>> print(r1.status, r1.reason)
200 OK
>>> data1 = r1.read() # This will return entire content.
>>> # The following example demonstrates reading data in chunks.
>>> conn.request("GET", "/")
>>> r1 = conn.getresponse()
>>> while not r1.closed:
...     print(r1.read(200)) # 200 bytes
b'<!doctype html>\n<!--[if"...
...
>>> # Example of an invalid request
>>> conn.request("GET", "/parrot.spam")
```



```

>>> r2 = conn.getresponse()
>>> print(r2.status, r2.reason)
404 Not Found
>>> data2 = r2.read()
>>> conn.close()

```

Here is an example session that uses the HEAD method. Note that the HEAD method never returns any data.

```

>>> import http.client
>>> conn = http.client.HTTPSConnection("www.python.org")
>>> conn.request("HEAD", "/")
>>> res = conn.getresponse()
>>> print(res.status, res.reason)
200 OK
>>> data = res.read()
>>> print(len(data))
0
>>> data == b''
True

```

Here is an example session that shows how to POST requests:

```

>>> import http.client, urllib.parse
>>> params = urllib.parse.urlencode({'@number': 12524, '@type': 'issue'})
>>> headers = {"Content-type": "application/x-www-form-urlencoded",
...           "Accept": "text/plain"}
>>> conn = http.client.HTTPConnection("bugs.python.org")
>>> conn.request("POST", "", params, headers)
>>> response = conn.getresponse()
>>> print(response.status, response.reason)
302 Found
>>> data = response.read()
>>> data
b'Redirecting to <a href="http://bugs.python.org/issue12524">http://bu
>>> conn.close()

```

Client side HTTP PUT requests are very similar to POST requests. The difference lies only the server side where HTTP server will allow resources to be created via PUT request. It should be noted that custom HTTP methods are also handled in `urllib.request.Request` by sending the appropriate `method` attribute. Here is an example session that shows how to do PUT request using `http.client`:

```

>>> # This creates an HTTP message
>>> # with the content of BODY as the enclosed representation
>>> # for the resource http://localhost:8080/file
...
>>> import http.client

```

```
>>> BODY = "***filecontents***"
>>> conn = http.client.HTTPConnection("localhost", 8080)
>>> conn.request("PUT", "/file", BODY)
>>> response = conn.getresponse()
>>> print(response.status, response.reason)
200, OK
```

## 21.12.4. HTTPMessage Objects

An `http.client.HTTPMessage` instance holds the headers from an HTTP response. It is implemented using the `email.message.Message` class.