

## 27.7. `tracemalloc` — Trace memory allocations

*New in version 3.4.*

**Source code:** [Lib/tracemalloc.py](#)

The `tracemalloc` module is a debug tool to trace memory blocks allocated by Python. It provides the following information:

- Traceback where an object was allocated
- Statistics on allocated memory blocks per filename and per line number: total size, number and average size of allocated memory blocks
- Compute the differences between two snapshots to detect memory leaks

To trace most memory blocks allocated by Python, the module should be started as early as possible by setting the `PYTHONTRACEMALLOC` environment variable to 1, or by using `-X tracemalloc` command line option. The `tracemalloc.start()` function can be called at runtime to start tracing Python memory allocations.

By default, a trace of an allocated memory block only stores the most recent frame (1 frame). To store 25 frames at startup: set the `PYTHONTRACEMALLOC` environment variable to 25, or use the `-X tracemalloc=25` command line option.

### 27.7.1. Examples

#### 27.7.1.1. Display the top 10

Display the 10 files allocating the most memory:

```
import tracemalloc

tracemalloc.start()

# ... run your application ...

snapshot = tracemalloc.take_snapshot()
top_stats = snapshot.statistics('lineno')

print("[ Top 10 ]")
for stat in top_stats[:10]:
    print(stat)
```

Example of output of the Python test suite:

```
[ Top 10 ]
<frozen importlib._bootstrap>:716: size=4855 KiB, count=39328, average=
<frozen importlib._bootstrap>:284: size=521 KiB, count=3199, average=1
/usr/lib/python3.4/collections/__init__.py:368: size=244 KiB, count=23
/usr/lib/python3.4/unittest/case.py:381: size=185 KiB, count=779, aver
/usr/lib/python3.4/unittest/case.py:402: size=154 KiB, count=378, aver
/usr/lib/python3.4/abc.py:133: size=88.7 KiB, count=347, average=262 B
<frozen importlib._bootstrap>:1446: size=70.4 KiB, count=911, average=
<frozen importlib._bootstrap>:1454: size=52.0 KiB, count=25, average=2
<string>:5: size=49.7 KiB, count=148, average=344 B
/usr/lib/python3.4/sysconfig.py:411: size=48.0 KiB, count=1, average=4
< >
```

We can see that Python loaded 4855 KiB data (bytecode and constants) from modules and that the `collections` module allocated 244 KiB to build `namedtuple` types.

See `Snapshot.statistics()` for more options.

## 27.7.1.2. Compute differences

Take two snapshots and display the differences:

```
import tracemalloc
tracemalloc.start()
# ... start your application ...

snapshot1 = tracemalloc.take_snapshot()
# ... call the function leaking memory ...
snapshot2 = tracemalloc.take_snapshot()

top_stats = snapshot2.compare_to(snapshot1, 'lineno')

print("[ Top 10 differences ]")
for stat in top_stats[:10]:
    print(stat)
```

Example of output before/after running some tests of the Python test suite:

```
[ Top 10 differences ]
<frozen importlib._bootstrap>:716: size=8173 KiB (+4428 KiB), count=71
/usr/lib/python3.4/linecache.py:127: size=940 KiB (+940 KiB), count=81
/usr/lib/python3.4/unittest/case.py:571: size=298 KiB (+298 KiB), cour
<frozen importlib._bootstrap>:284: size=1005 KiB (+166 KiB), count=742
/usr/lib/python3.4/mimetypes.py:217: size=112 KiB (+112 KiB), count=13
/usr/lib/python3.4/http/server.py:848: size=96.0 KiB (+96.0 KiB), cour
/usr/lib/python3.4/inspect.py:1465: size=83.5 KiB (+83.5 KiB), count=1
/usr/lib/python3.4/unittest/mock.py:491: size=77.7 KiB (+77.7 KiB), cc
```

```
/usr/lib/python3.4/urllib/parse.py:476: size=71.8 KiB (+71.8 KiB), cou  
/usr/lib/python3.4/contextlib.py:38: size=67.2 KiB (+67.2 KiB), count=
```

We can see that Python has loaded 8173 KiB of module data (bytecode and constants), and that this is 4428 KiB more than had been loaded before the tests, when the previous snapshot was taken. Similarly, the `linecache` module has cached 940 KiB of Python source code to format tracebacks, all of it since the previous snapshot.

If the system has little free memory, snapshots can be written on disk using the `Snapshot.dump()` method to analyze the snapshot offline. Then use the `Snapshot.load()` method reload the snapshot.

### 27.7.1.3. Get the traceback of a memory block

Code to display the traceback of the biggest memory block:

```
import tracemalloc

# Store 25 frames
tracemalloc.start(25)

# ... run your application ...

snapshot = tracemalloc.take_snapshot()
top_stats = snapshot.statistics('traceback')

# pick the biggest memory block
stat = top_stats[0]
print("%s memory blocks: %.1f KiB" % (stat.count, stat.size / 1024))
for line in stat.traceback.format():
    print(line)
```

Example of output of the Python test suite (traceback limited to 25 frames):

```
903 memory blocks: 870.1 KiB
File "<frozen importlib._bootstrap>", line 716
File "<frozen importlib._bootstrap>", line 1036
File "<frozen importlib._bootstrap>", line 934
File "<frozen importlib._bootstrap>", line 1068
File "<frozen importlib._bootstrap>", line 619
File "<frozen importlib._bootstrap>", line 1581
File "<frozen importlib._bootstrap>", line 1614
File "/usr/lib/python3.4/doctest.py", line 101
import pdb
File "<frozen importlib._bootstrap>", line 284
File "<frozen importlib._bootstrap>", line 938
File "<frozen importlib._bootstrap>", line 1068
```

```

File "<frozen importlib._bootstrap>", line 619
File "<frozen importlib._bootstrap>", line 1581
File "<frozen importlib._bootstrap>", line 1614
File "/usr/lib/python3.4/test/support/__init__.py", line 1728
    import doctest
File "/usr/lib/python3.4/test/test_pickletools.py", line 21
    support.run_doctest(pickletools)
File "/usr/lib/python3.4/test/regrtest.py", line 1276
    test_runner()
File "/usr/lib/python3.4/test/regrtest.py", line 976
    display_failure=not verbose)
File "/usr/lib/python3.4/test/regrtest.py", line 761
    match_tests=ns.match_tests)
File "/usr/lib/python3.4/test/regrtest.py", line 1563
    main()
File "/usr/lib/python3.4/test/__main__.py", line 3
    regrtest.main_in_temp_cwd()
File "/usr/lib/python3.4/runpy.py", line 73
    exec(code, run_globals)
File "/usr/lib/python3.4/runpy.py", line 160
    "__main__", fname, loader, pkg_name)

```

We can see that the most memory was allocated in the `importlib` module to load data (bytecode and constants) from modules: 870.1 KiB. The traceback is where the `importlib` loaded data most recently: on the `import pdb` line of the `doctest` module. The traceback may change if a new module is loaded.

## 27.7.1.4. Pretty top

Code to display the 10 lines allocating the most memory with a pretty output, ignoring `<frozen importlib._bootstrap>` and `<unknown>` files:

```

import linecache
import os
import tracemalloc

def display_top(snapshot, key_type='lineno', limit=10):
    snapshot = snapshot.filter_traces((
        tracemalloc.Filter(False, "<frozen importlib._bootstrap>"),
        tracemalloc.Filter(False, "<unknown>"),
    ))
    top_stats = snapshot.statistics(key_type)

    print("Top %s lines" % limit)
    for index, stat in enumerate(top_stats[:limit], 1):
        frame = stat.traceback[0]
        # replace "/path/to/module/file.py" with "module/file.py"
        filename = os.sep.join(frame.filename.split(os.sep)[-2:])
        print("#%s: %s:%s: %.1f KiB"

```

```

        % (index, filename, frame.lineno, stat.size / 1024))
    line = linecache.getline(frame.filename, frame.lineno).strip()
    if line:
        print('    %s' % line)

    other = top_stats[limit:]
    if other:
        size = sum(stat.size for stat in other)
        print("%s other: %.1f KiB" % (len(other), size / 1024))
    total = sum(stat.size for stat in top_stats)
    print("Total allocated size: %.1f KiB" % (total / 1024))

tracemalloc.start()

# ... run your application ...

snapshot = tracemalloc.take_snapshot()
display_top(snapshot)

```

Example of output of the Python test suite:

```

Top 10 lines
#1: Lib/base64.py:414: 419.8 KiB
    _b85chars2 = [(a + b) for a in _b85chars for b in _b85chars]
#2: Lib/base64.py:306: 419.8 KiB
    _a85chars2 = [(a + b) for a in _a85chars for b in _a85chars]
#3: collections/__init__.py:368: 293.6 KiB
    exec(class_definition, namespace)
#4: Lib/abc.py:133: 115.2 KiB
    cls = super().__new__(mcls, name, bases, namespace)
#5: unittest/case.py:574: 103.1 KiB
    testMethod()
#6: Lib/linecache.py:127: 95.4 KiB
    lines = fp.readlines()
#7: urllib/parse.py:476: 71.8 KiB
    for a in _hexdig for b in _hexdig}
#8: <string>:5: 62.0 KiB
#9: Lib/_weakrefset.py:37: 60.0 KiB
    self.data = set()
#10: Lib/base64.py:142: 59.8 KiB
    _b32tab2 = [a + b for a in _b32tab for b in _b32tab]
6220 other: 3602.8 KiB
Total allocated size: 5303.1 KiB

```

See [Snapshot.statistics\(\)](#) for more options.

## 27.7.2. API

### 27.7.2.1. Functions

`tracemalloc.clear_traces()`

Clear traces of memory blocks allocated by Python.

See also [stop\(\)](#).

`tracemalloc.get_object_traceback(obj)`

Get the traceback where the Python object *obj* was allocated. Return a [Traceback](#) instance, or None if the [tracemalloc](#) module is not tracing memory allocations or did not trace the allocation of the object.

See also [gc.get\\_referrers\(\)](#) and [sys.getsizeof\(\)](#) functions.

`tracemalloc.get_traceback_limit()`

Get the maximum number of frames stored in the traceback of a trace.

The [tracemalloc](#) module must be tracing memory allocations to get the limit, otherwise an exception is raised.

The limit is set by the [start\(\)](#) function.

`tracemalloc.get_traced_memory()`

Get the current size and peak size of memory blocks traced by the [tracemalloc](#) module as a tuple: (current: int, peak: int).

`tracemalloc.get_tracemalloc_memory()`

Get the memory usage in bytes of the [tracemalloc](#) module used to store traces of memory blocks. Return an [int](#).

`tracemalloc.is_tracing()`

True if the [tracemalloc](#) module is tracing Python memory allocations, False otherwise.

See also [start\(\)](#) and [stop\(\)](#) functions.

`tracemalloc.start(nframe: int=1)`

Start tracing Python memory allocations: install hooks on Python memory allocators. Collected tracebacks of traces will be limited to *nframe* frames. By default, a trace of a memory block only stores the most recent frame: the limit is 1. *nframe* must be greater or equal to 1.

Storing more than 1 frame is only useful to compute statistics grouped by 'traceback' or to compute cumulative statistics: see the [Snapshot.compare\\_to\(\)](#) and [Snapshot.statistics\(\)](#) methods.

Storing more frames increases the memory and CPU overhead of the [tracemalloc](#) module. Use the [get\\_tracemalloc\\_memory\(\)](#) function to measure how much memory is used by the [tracemalloc](#) module.

The [PYTHONTRACEMALLOC](#) environment variable ([PYTHONTRACEMALLOC=NFRAME](#)) and the [-X tracemalloc=NFRAME](#) command line option can be used to start tracing at startup.

See also [stop\(\)](#), [is\\_tracing\(\)](#) and [get\\_traceback\\_limit\(\)](#) functions.

#### `tracemalloc.stop()`

Stop tracing Python memory allocations: uninstall hooks on Python memory allocators. Also clears all previously collected traces of memory blocks allocated by Python.

Call [take\\_snapshot\(\)](#) function to take a snapshot of traces before clearing them.

See also [start\(\)](#), [is\\_tracing\(\)](#) and [clear\\_traces\(\)](#) functions.

#### `tracemalloc.take_snapshot()`

Take a snapshot of traces of memory blocks allocated by Python. Return a new [Snapshot](#) instance.

The snapshot does not include memory blocks allocated before the [tracemalloc](#) module started to trace memory allocations.

Tracebacks of traces are limited to [get\\_traceback\\_limit\(\)](#) frames. Use the *nframe* parameter of the [start\(\)](#) function to store more frames.

The [tracemalloc](#) module must be tracing memory allocations to take a snapshot, see the [start\(\)](#) function.

See also the [get\\_object\\_traceback\(\)](#) function.

### 27.7.2.2. DomainFilter

*class* `tracemalloc.DomainFilter`(*inclusive: bool, domain: int*)

Filter traces of memory blocks by their address space (domain).

*New in version 3.6.*

## **inclusive**

If *inclusive* is True (include), match memory blocks allocated in the address space [domain](#).

If *inclusive* is False (exclude), match memory blocks not allocated in the address space [domain](#).

## **domain**

Address space of a memory block (int). Read-only property.

### 27.7.2.3. Filter

```
class tracemalloc.Filter(inclusive: bool, filename_pattern: str, lineno:
int=None, all_frames: bool=False, domain: int=None)
```

Filter on traces of memory blocks.

See the [fnmatch.fnmatch\(\)](#) function for the syntax of *filename\_pattern*. The '.pyc' file extension is replaced with '.py'.

Examples:

- `Filter(True, subprocess.__file__)` only includes traces of the [subprocess](#) module
- `Filter(False, tracemalloc.__file__)` excludes traces of the [tracemalloc](#) module
- `Filter(False, "<unknown>")` excludes empty tracebacks

*Changed in version 3.5:* The '.pyo' file extension is no longer replaced with '.py'.

*Changed in version 3.6:* Added the [domain](#) attribute.

## **domain**

Address space of a memory block (int or None).

## **inclusive**

If *inclusive* is True (include), only match memory blocks allocated in a file with a name matching [filename\\_pattern](#) at line number [lineno](#).

If *inclusive* is False (exclude), ignore memory blocks allocated in a file with a name matching [filename\\_pattern](#) at line number [lineno](#).

## **lineno**

Line number (int) of the filter. If *lineno* is None, the filter matches any line number.



## **filename\_pattern**

Filename pattern of the filter (str). Read-only property.

## **all\_frames**

If *all\_frames* is True, all frames of the traceback are checked. If *all\_frames* is False, only the most recent frame is checked.

This attribute has no effect if the traceback limit is 1. See the [get\\_traceback\\_limit\(\)](#) function and [Snapshot.traceback\\_limit](#) attribute.

## 27.7.2.4. Frame

*class* tracemalloc.**Frame**

Frame of a traceback.

The [Traceback](#) class is a sequence of [Frame](#) instances.

### **filename**

Filename (str).

### **lineno**

Line number (int).

## 27.7.2.5. Snapshot

*class* tracemalloc.**Snapshot**

Snapshot of traces of memory blocks allocated by Python.

The [take\\_snapshot\(\)](#) function creates a snapshot instance.

**compare\_to**(*old\_snapshot: Snapshot, key\_type: str, cumulative: bool=False*)

Compute the differences with an old snapshot. Get statistics as a sorted list of [StatisticDiff](#) instances grouped by *key\_type*.

See the [Snapshot.statistics\(\)](#) method for *key\_type* and *cumulative* parameters.

The result is sorted from the biggest to the smallest by: absolute value of [StatisticDiff.size\\_diff](#), [StatisticDiff.size](#), absolute value of [StatisticDiff.count\\_diff](#), [StatisticDiff.count](#) and then by [StatisticDiff.traceback](#).

## **dump(filename)**

Write the snapshot into a file.

Use [load\(\)](#) to reload the snapshot.

## **filter\_traces(filters)**

Create a new [Snapshot](#) instance with a filtered [traces](#) sequence, *filters* is a list of [DomainFilter](#) and [Filter](#) instances. If *filters* is an empty list, return a new [Snapshot](#) instance with a copy of the traces.

All inclusive filters are applied at once, a trace is ignored if no inclusive filters match it. A trace is ignored if at least one exclusive filter matches it.

*Changed in version 3.6:* [DomainFilter](#) instances are now also accepted in *filters*.

## *classmethod* **load(filename)**

Load a snapshot from a file.

See also [dump\(\)](#).

## **statistics(key\_type: str, cumulative: bool=False)**

Get statistics as a sorted list of [Statistic](#) instances grouped by *key\_type*:

key_type	description
'filename'	filename
'lineno'	filename and line number
'traceback'	traceback

If *cumulative* is `True`, cumulate size and count of memory blocks of all frames of the traceback of a trace, not only the most recent frame. The cumulative mode can only be used with *key\_type* equals to 'filename' and 'lineno'.

The result is sorted from the biggest to the smallest by: [Statistic.size](#), [Statistic.count](#) and then by [Statistic.traceback](#).

## **traceback\_limit**

Maximum number of frames stored in the traceback of [traces](#): result of the [get\\_traceback\\_limit\(\)](#) when the snapshot was taken.

## **traces**

Traces of all memory blocks allocated by Python: sequence of [Trace](#) instances.

The sequence has an undefined order. Use the [Snapshot.statistics\(\)](#) method to get a sorted list of statistics.

### 27.7.2.6. Statistic

*class* tracemalloc.**Statistic**

Statistic on memory allocations.

[Snapshot.statistics\(\)](#) returns a list of [Statistic](#) instances.

See also the [StatisticDiff](#) class.

**count**

Number of memory blocks ([int](#)).

**size**

Total size of memory blocks in bytes ([int](#)).

**traceback**

Traceback where the memory block was allocated, [Traceback](#) instance.

### 27.7.2.7. StatisticDiff

*class* tracemalloc.**StatisticDiff**

Statistic difference on memory allocations between an old and a new [Snapshot](#) instance.

[Snapshot.compare\\_to\(\)](#) returns a list of [StatisticDiff](#) instances. See also the [Statistic](#) class.

**count**

Number of memory blocks in the new snapshot ([int](#)): 0 if the memory blocks have been released in the new snapshot.

**count\_diff**

Difference of number of memory blocks between the old and the new snapshots ([int](#)): 0 if the memory blocks have been allocated in the new snapshot.

**size**

Total size of memory blocks in bytes in the new snapshot ([int](#)): 0 if the memory blocks have been released in the new snapshot.

**size\_diff**

Difference of total size of memory blocks in bytes between the old and the new snapshots (int): 0 if the memory blocks have been allocated in the new snapshot.

### **traceback**

Traceback where the memory blocks were allocated, [Traceback](#) instance.

## 27.7.2.8. Trace

### *class* tracemalloc.**Trace**

Trace of a memory block.

The [Snapshot.traces](#) attribute is a sequence of [Trace](#) instances.

### **size**

Size of the memory block in bytes (int).

### **traceback**

Traceback where the memory block was allocated, [Traceback](#) instance.

## 27.7.2.9. Traceback

### *class* tracemalloc.**Traceback**

Sequence of [Frame](#) instances sorted from the most recent frame to the oldest frame.

A traceback contains at least 1 frame. If the tracemalloc module failed to get a frame, the filename "<unknown>" at line number 0 is used.

When a snapshot is taken, tracebacks of traces are limited to [get\\_traceback\\_limit\(\)](#) frames. See the [take\\_snapshot\(\)](#) function.

The [Trace.traceback](#) attribute is an instance of [Traceback](#) instance.

### **format**(*limit=None*)

Format the traceback as a list of lines with newlines. Use the [linecache](#) module to retrieve lines from the source code. If *limit* is set, only format the *limit* most recent frames.

Similar to the [traceback.format\\_tb\(\)](#) function, except that [format\(\)](#) does not include newlines.

Example:

```
print("Traceback (most recent call first):")
for line in traceback:
    print(line)
```

Output:

```
Traceback (most recent call first):
  File "test.py", line 9
    obj = Object()
  File "test.py", line 12
    tb = tracemalloc.get_object_traceback(f())
```